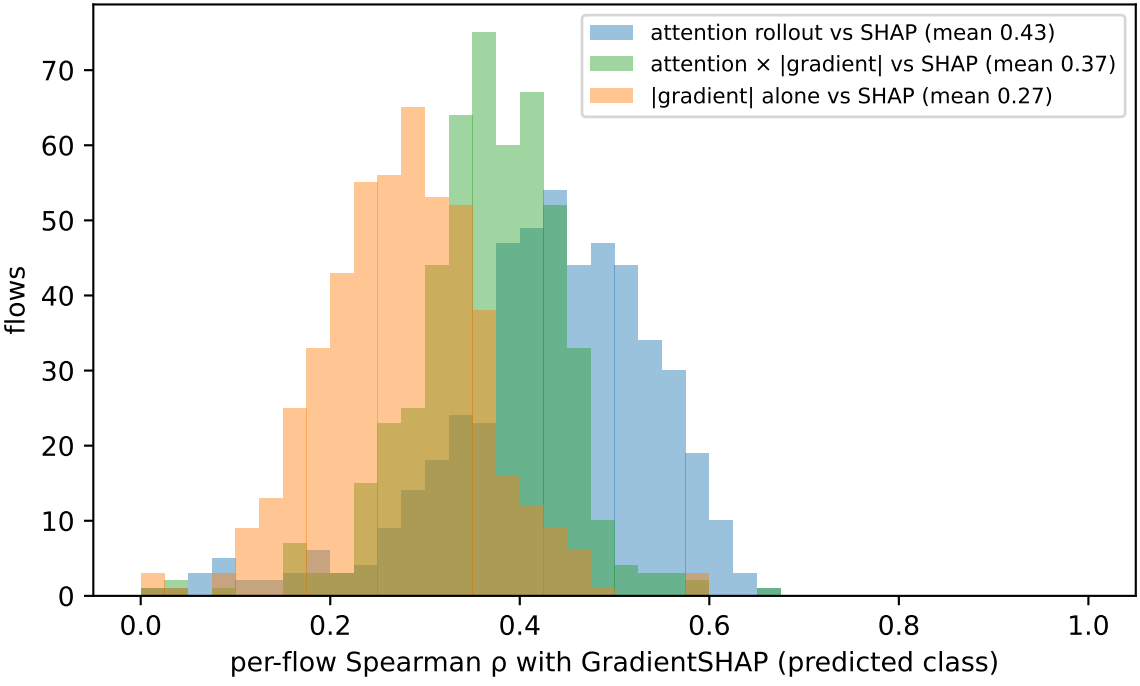
Consistency between built-in and post-hoc attributions (500 flows)
attention rollout vs SHAP (mean 0.43)
attention × |gradient| vs SHAP (mean 0.37)
|gradient| alone vs SHAP (mean 0.27)
flows
0
10
20
30
40
50
60
70
0.0
0.2
0.4
0.6
0.8
1.0
per-flow Spearman ρ with GradientSHAP (predicted class)